
AES-128 ECB

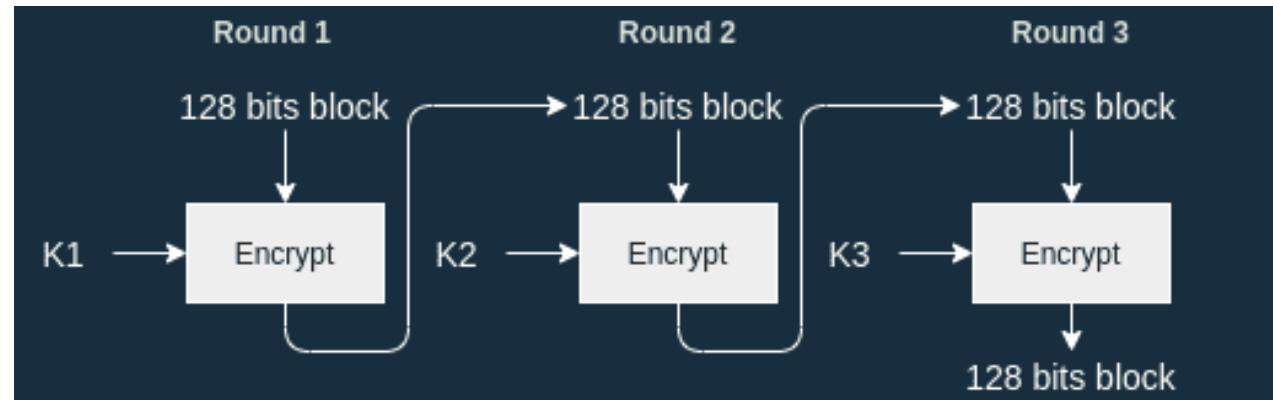
Comparison between sequential,
OpenMP and CUDA
implementations

Daniele Bianchini
HPC Assignment

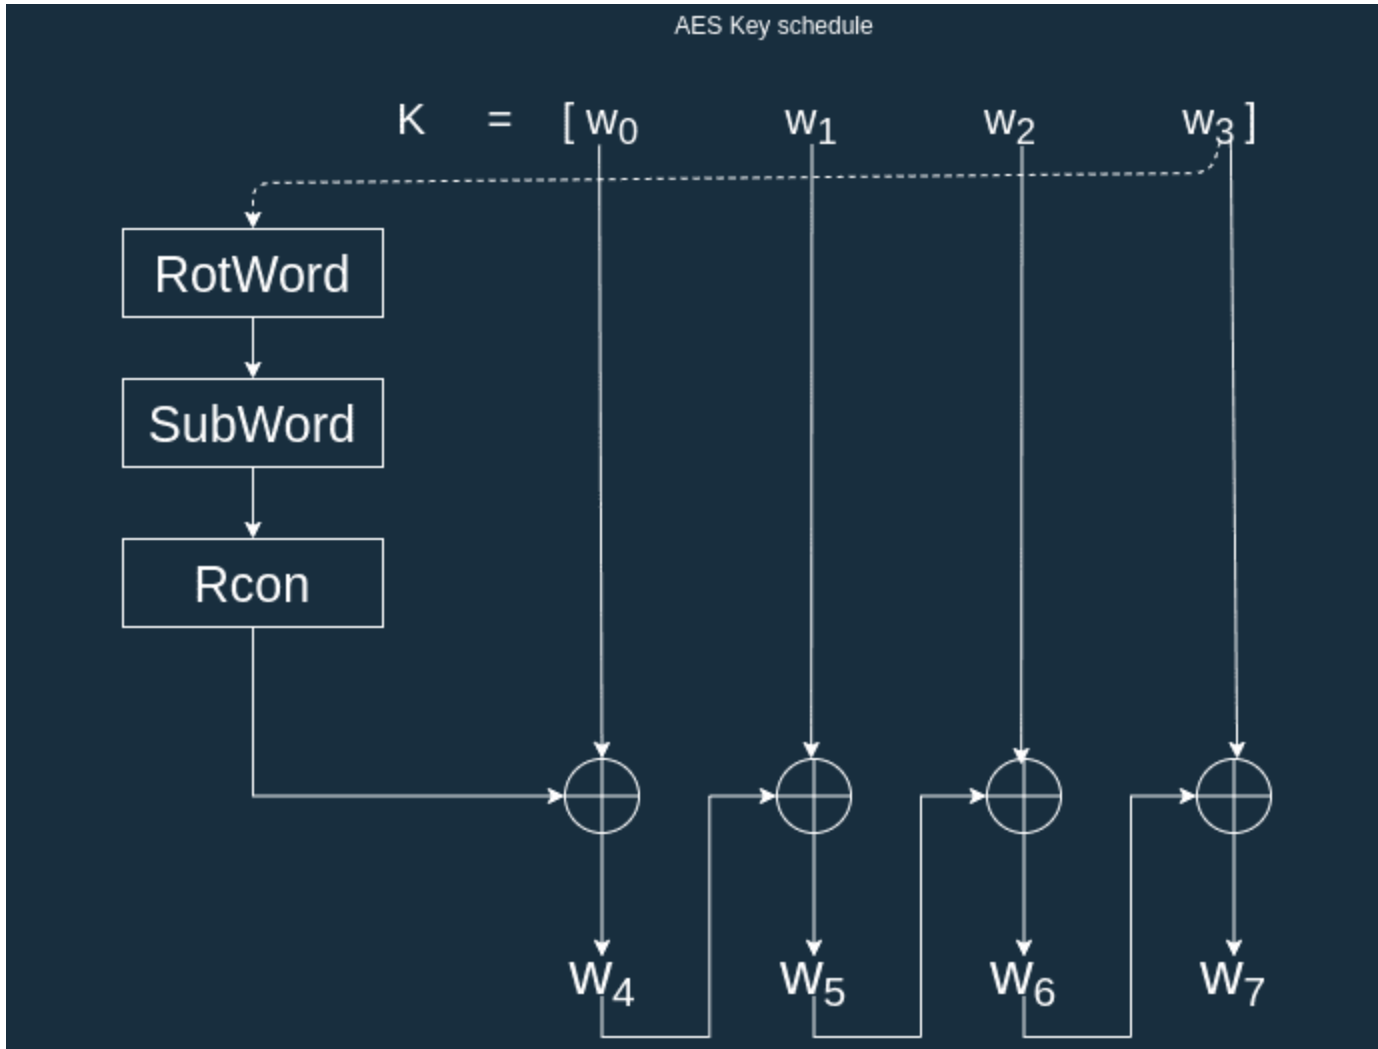


AES-128

- Plaintext divided in 16 bytes blocks
- 16 byte long secret key (128 bits)
- 10 round keys derived from secret key -> Key scheduling
- Electronic codebook mode of operation: each block is ciphered independently
- Each block undergoes 11 rounds with each roundkey
- S-Box: constant matrix used to permute bytes



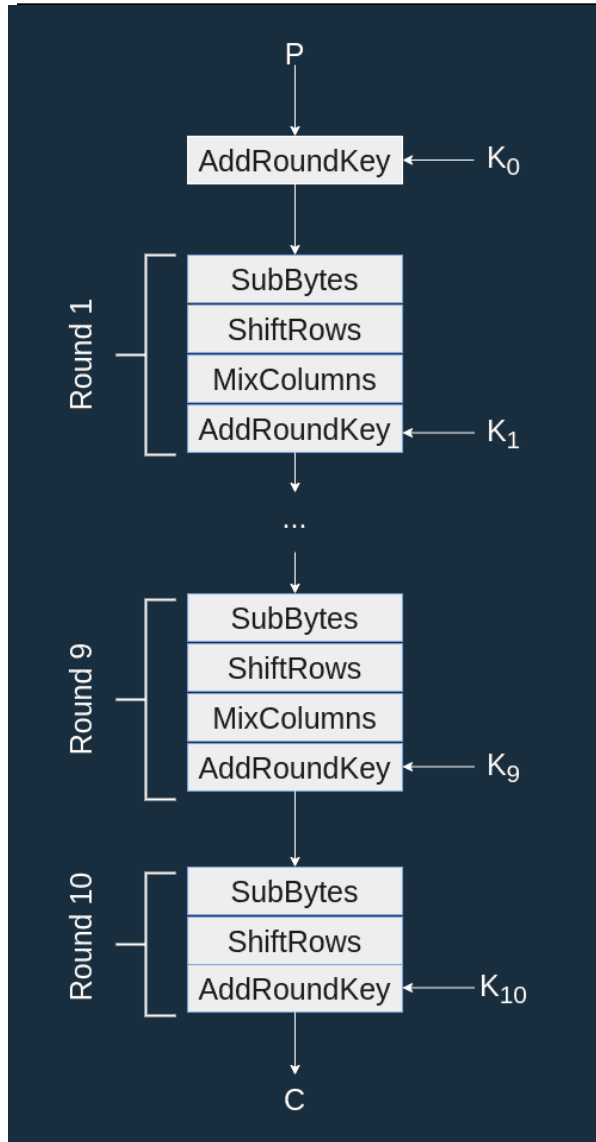
Key Scheduling



- RotWord: Left-rotate a 4 byte word
 - $(w_0\ w_1\ w_2\ w_3) \rightarrow (w_1\ w_2\ w_3\ w_0)$
- SubWord: Apply S-Box permutation to each byte of a word
- Rcon: xor first byte of a word with a constant round r dependant $\rightarrow 2^{r-1} \bmod 2^8$

Until we obtain 10 keys. Dependencies prevents parallelization

Cipher algorithm



- **AddRoundKey**: xor block of plaintext with round key
- **ShiftRows**: performs a cyclical rotation of the block by rows of bytes (matrix of the state)


$$\begin{bmatrix} s_{0,0} & s_{0,1} & s_{0,2} & s_{0,3} \\ s_{1,0} & s_{1,1} & s_{1,2} & s_{1,3} \\ s_{2,0} & s_{2,1} & s_{2,2} & s_{2,3} \\ s_{3,0} & s_{3,1} & s_{3,2} & s_{3,3} \end{bmatrix} \rightarrow \begin{bmatrix} s_{0,0} & s_{0,1} & s_{0,2} & s_{0,3} \\ s_{1,1} & s_{1,2} & s_{1,3} & s_{1,0} \\ s_{2,2} & s_{2,3} & s_{2,0} & s_{2,1} \\ s_{3,3} & s_{3,0} & s_{3,1} & s_{3,2} \end{bmatrix}$$

- **MixColumns**: Matrix constant multiplication in $GF\ 2^8$

$$\begin{bmatrix} t_{0,0} \\ t_{1,0} \\ t_{2,0} \\ t_{3,0} \end{bmatrix} = \begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} s_{0,0} \\ s_{1,0} \\ s_{2,0} \\ s_{3,0} \end{bmatrix}$$

Improving performance

- Bigger plaintext means more blocks to cipher, longer time of execution
- Key scheduling is relatively fast and not parallelizable
- Cipher function treats each block independently
 - synchronization ok
 - Static workload
- Cipher function treats contiguous addresses of blocks
- Cipher function can be unrolled very well



Electronic Codebook (ECB) mode encryption

OpenMP

- Max improvement with 16 threads
- For cycle iterating on each block
- Static and balanced workload

```
omp_set_num_threads(16);

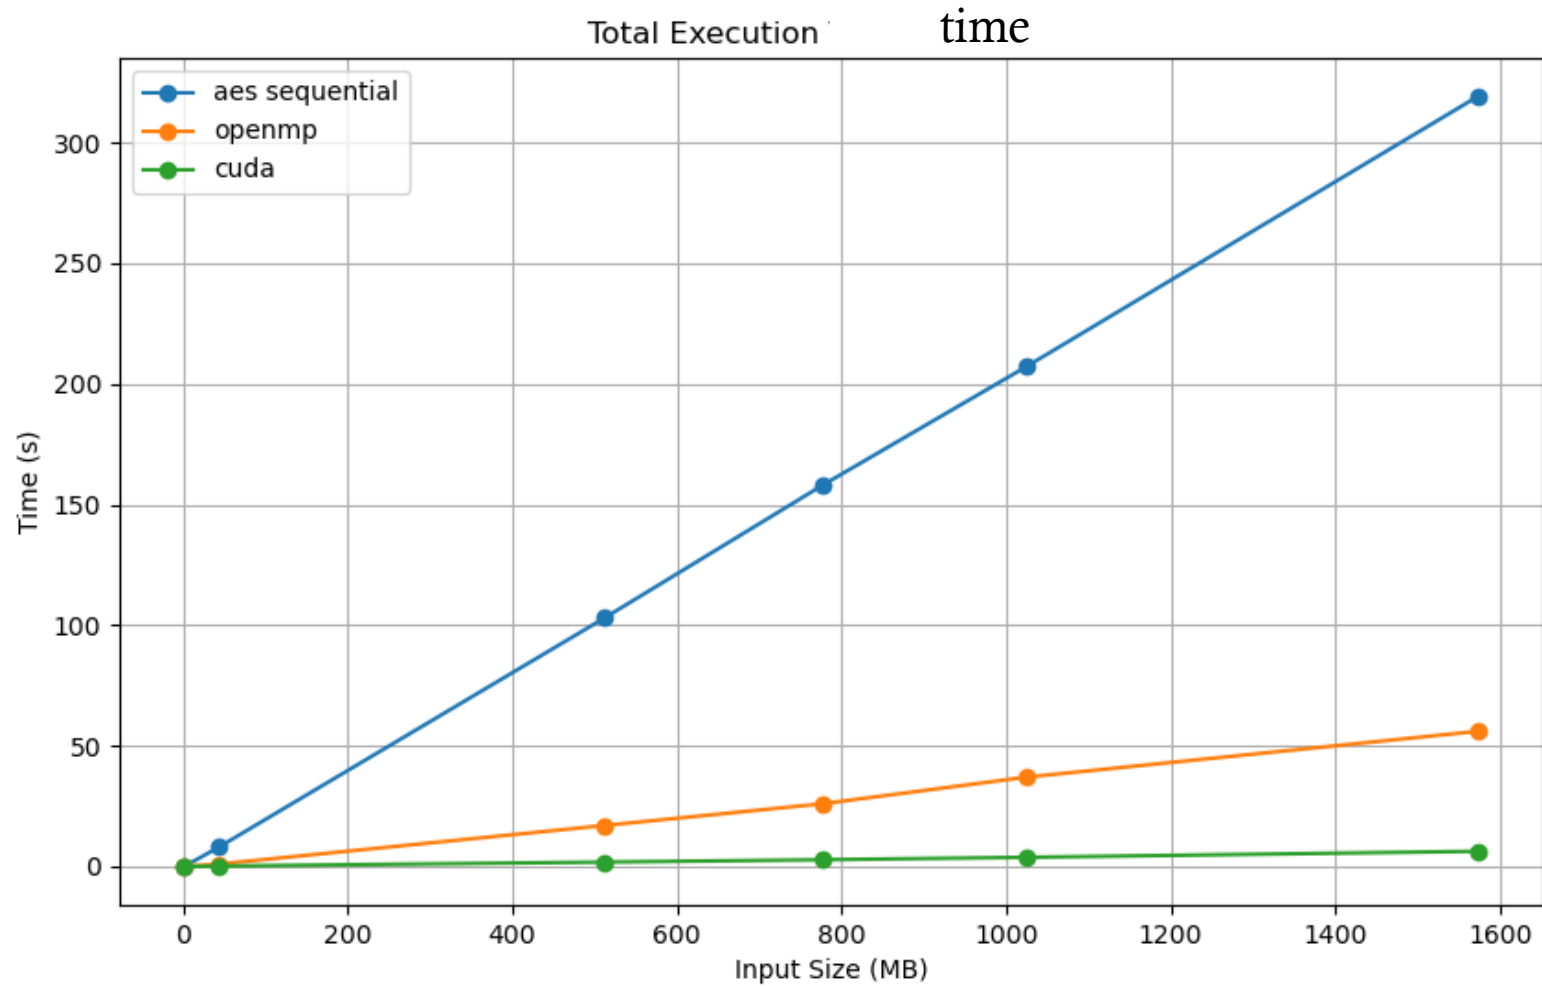
printf("Encrypting...\n");
#pragma omp parallel for
for(int i = 0; i < blocks_to_cipher; i++){
    encrypt_ECB(plaintext + (i*BLOCKSIZE), KEY);
}
```

CUDA

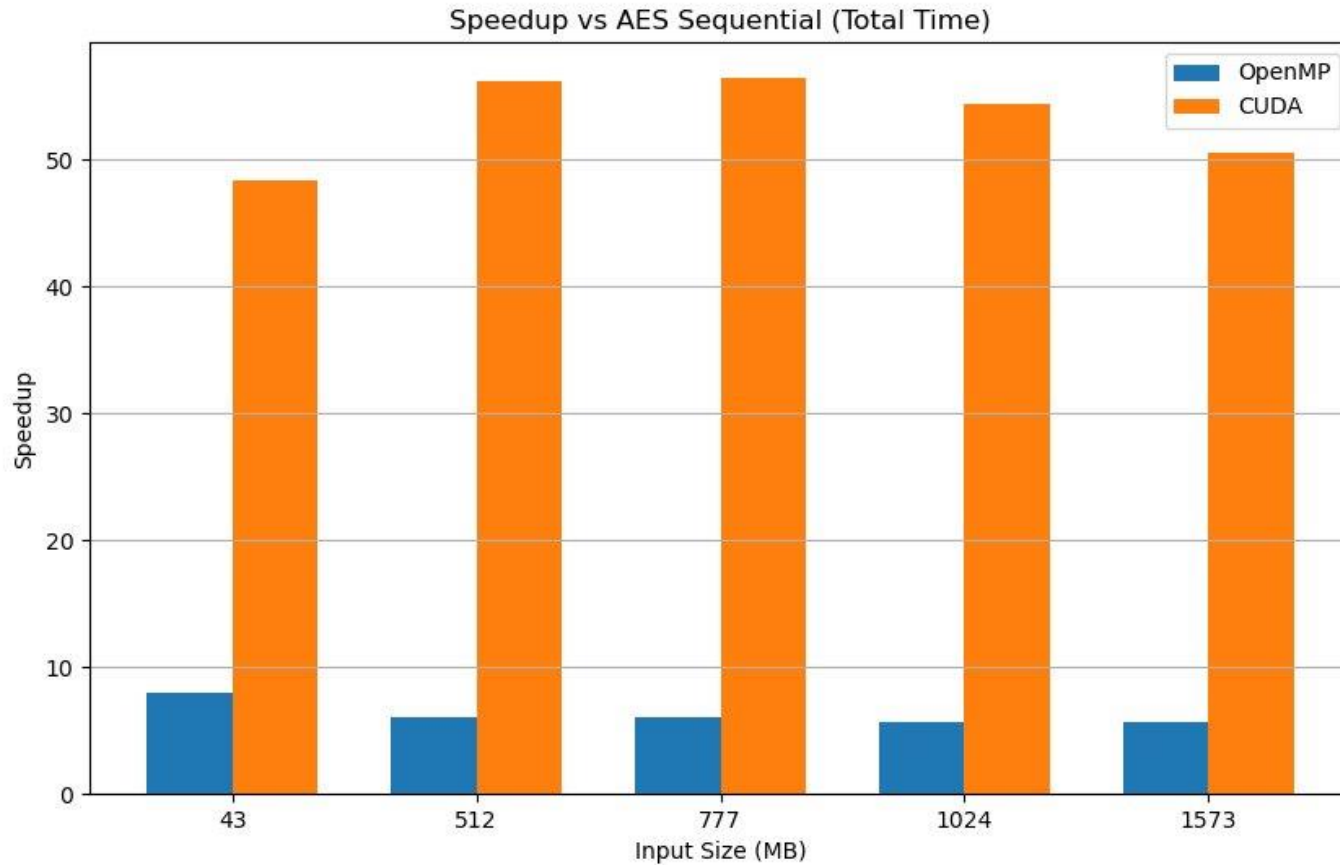
- S-Box and keys in constant memory for fast read-only access
- Optimal number of threads: 256
- Cuda blocks: $(\text{blocks_to_cipher} + \text{thread_n} - 1) / \text{thread_n}$
- Each thread works on the block corresponding to its ID and all blocks at offset $(\text{blockDim.x} * \text{gridDim.x})$

```
int thread_n = 256;
int blocks_pergrid = (blocks_to_cipher + thread_n - 1) / thread_n;
printf("Blocks: %d\n", blocks_pergrid);
cudaEventRecord(alg_start);
encryptFull_ECB<<<blocks_pergrid, thread_n>>>(plain_d, blocks_to_cipher);
if (err != cudaSuccess) {
    printf("CUDA kernel launch failed: %s\n", cudaGetErrorString(err));
}
cudaDeviceSynchronize();
cudaEventRecord(alg_stop);
```

Execution time




Speedup



- OpenMP speedup is about 5-7x
- Cuda speedup up to >50x
- With bigger filesizes to cipher, speedup lowers for CUDA...

Non-kernel execution time



- Bigger plaintexts means more memory to copy to and from the device
- Non-kernel execution time grows fast
- Speedup compensate more than well

Github Repo

- `git clone https://github.com/Lick1Fonzi/AES-128_ECB_Parallel_HPCUnimore.git`
 - `./xstat.sh` -> generates 4 pngs with the result charts
-